

SHRIJAL GOSWAMI

✉ goswamivansh999@gmail.com · [in linkedin.com/in/shrijalgoswami](https://www.linkedin.com/in/shrijalgoswami) · github.com/Shri-AI-ML · 📍 Bhopal, India

SUMMARY

Second-year B.Tech AI/ML student at VIT Bhopal with production experience building RAG pipelines, stacking ensembles, and LLM-integrated REST backends. Deployed end-to-end ML systems achieving **89.13% accuracy** / **0.94 AUC**; currently interning at Pinnacle Labs. Seeking AI/ML & Data Science internships for Summer/Fall 2026.

EDUCATION

B.Tech — Computer Science & Engineering (AI/ML) | VIT Bhopal University

Sep 2024 – Oct 2028

TECHNICAL SKILLS

Languages: Python, C++, SQL

ML / AI: Scikit-learn, XGBoost, Stacking Ensembles, BM25, Sentence-Transformers, PyTorch (fundamentals)

GenAI / LLMs: LangChain, RAG Pipelines, ChromaDB, Groq API, Gemini, Vertex AI, Prompt Engineering, Explainable AI

Frameworks & Data: FastAPI, Streamlit, NumPy, Pandas, Matplotlib, Seaborn, Feature Engineering, EDA

Cloud / Tools: Google Cloud (Vertex AI), Git/GitHub, Jupyter Notebook, VS Code

EXPERIENCE

Artificial Intelligence Intern | Pinnacle Labs

May 2026 – Present | Kolkata (Remote)

- Built modular Python AI components integrating applied ML workflows with production-grade FastAPI backends across iterative development cycles.
- Enhanced and debugged ML pipelines — improving code reliability, feature correctness, and deployment readiness for AI-focused deliverables.

PROJECTS

Explainable RAG QA System

github.com/Shri-AI-ML/Explainable-RAG-QA-System

Python · FastAPI · ChromaDB · Groq LLM · BM25 · LangChain · Hugging Face Embeddings

- Architected hybrid retrieval pipeline (BM25 sparse + dense vector search) over ChromaDB with Groq LLM backend, delivering grounded, source-attributed answers with explainability.
- Built incremental document indexing eliminating full DB rebuilds on new ingestion — supports PDF/JSON with automated chunking and Hugging Face embedding inference optimized for low-memory deployment.
- Exposed via FastAPI REST endpoints; tuned hybrid re-ranking weights and chunk sizing to minimize retrieval latency at production scale.

Heart Disease Prediction System

github.com/Shri-AI-ML/Heart-Disease-Prediction

Python · Scikit-learn · XGBoost · FastAPI · Streamlit · ChromaDB · LangChain

- Engineered hybrid stacking ensemble (Random Forest + XGBoost + Logistic Regression) on UCI Heart Disease dataset — **89.13% accuracy** and **0.94 AUC** via cross-validated stacked generalization.
- Executed full ML pipeline: EDA → outlier detection → StandardScaler normalization → feature selection → stacked cross-validation for reproducibility and clinical interpretability.
- Deployed end-to-end: FastAPI prediction REST API + Streamlit frontend; integrated ChromaDB + LangChain RAG layer for clinical Q&A on patient data.

CERTIFICATIONS & ACHIEVEMENTS

- **DataForge Hackathon** — Finalist, IITR E-Summit '26, VIT Bhopal
- Applied Machine Learning in Python — *Coursera / University of Michigan*
- Build Real-World AI Apps with Gemini & Imagen — *Google Cloud Skills Boost*
- Generative AI Apps with Gemini & Streamlit — *Google Cloud Skills Boost*
- Inspect Rich Documents with Gemini Multimodality & Multimodal RAG — *Google Cloud*